

CSE4088 Introduction to Machine Learning

Homework 3

Berk Kirtay - 150118043

Gradient Descent

For this section, Gradient_Descent.py is implemented. Script can be run to observe results. Plotting is included.

Results after running Gradient_Descent.py:

```
$ py Gradient_Descent.py
Gradient Descent:
Number of iterations: 10
Final(u,v): (0.04473629039778207, 0.023958714099141746)
Error rate after completion: 1.2086833944220747e-15
Coordinate Descent:
Number of iterations: 15
Final(u,v): (6.29707589930517, -2.852306954077811)
Error rate after completion: 0.13981379199615315
```

4.

The partial derivative of $E(u,v)$ with respect to $\frac{\partial E}{\partial u}$ is

$$\frac{\partial E}{\partial u} = 2 \cdot (u \cdot e^v - 2 \cdot v \cdot e^{-u}) \cdot (e^v + 2 \cdot v \cdot e^{-u})$$

5. and 6.

In this section we apply gradient descent algorithm to the nonlinear error space function by calculating the derivatives of squared residuals until error becomes less than 10^{-14} . Derivatives to apply gradient descent:


$$\frac{\partial E}{\partial u} = 2 \cdot (u \cdot e^v - 2 \cdot v \cdot e^{-u}) \cdot (e^v + 2 \cdot v \cdot e^{-u})$$

$$\frac{\partial E}{\partial v} = 2 \cdot (u \cdot e^v - 2 \cdot v \cdot e^{-u}) \cdot (u \cdot e^v - 2 \cdot e^{-u})$$

Number of iterations until error becomes $1.2086833944220747e-15$ is 10.

The latest (u,v) pair is $(0.04473629039778207, 0.023958714099141746)$

Plotting of the (u,v) pair on every iteration:




The answers are 5: D, and 6: E

7.

The error does not converge and stops at 0.13981379199615315 which is close to 10^{-1}

The latest (u,v) pair is $(6.29707589930517, -2.852306954077811)$

Plotting of the (u,v) pair on every iteration:



The answer is A.

Logistic Regression

For this section, Logistic_Regression.py is implemented. Script can be run to observe results. Plotting is included.

Results after running Logistic_Regression.py:

```
$ py Logistic_Regression.py
Eout: 0.09444365341119473, epochs: 338.58
```

The used formulas for the logistic regression algorithm are:

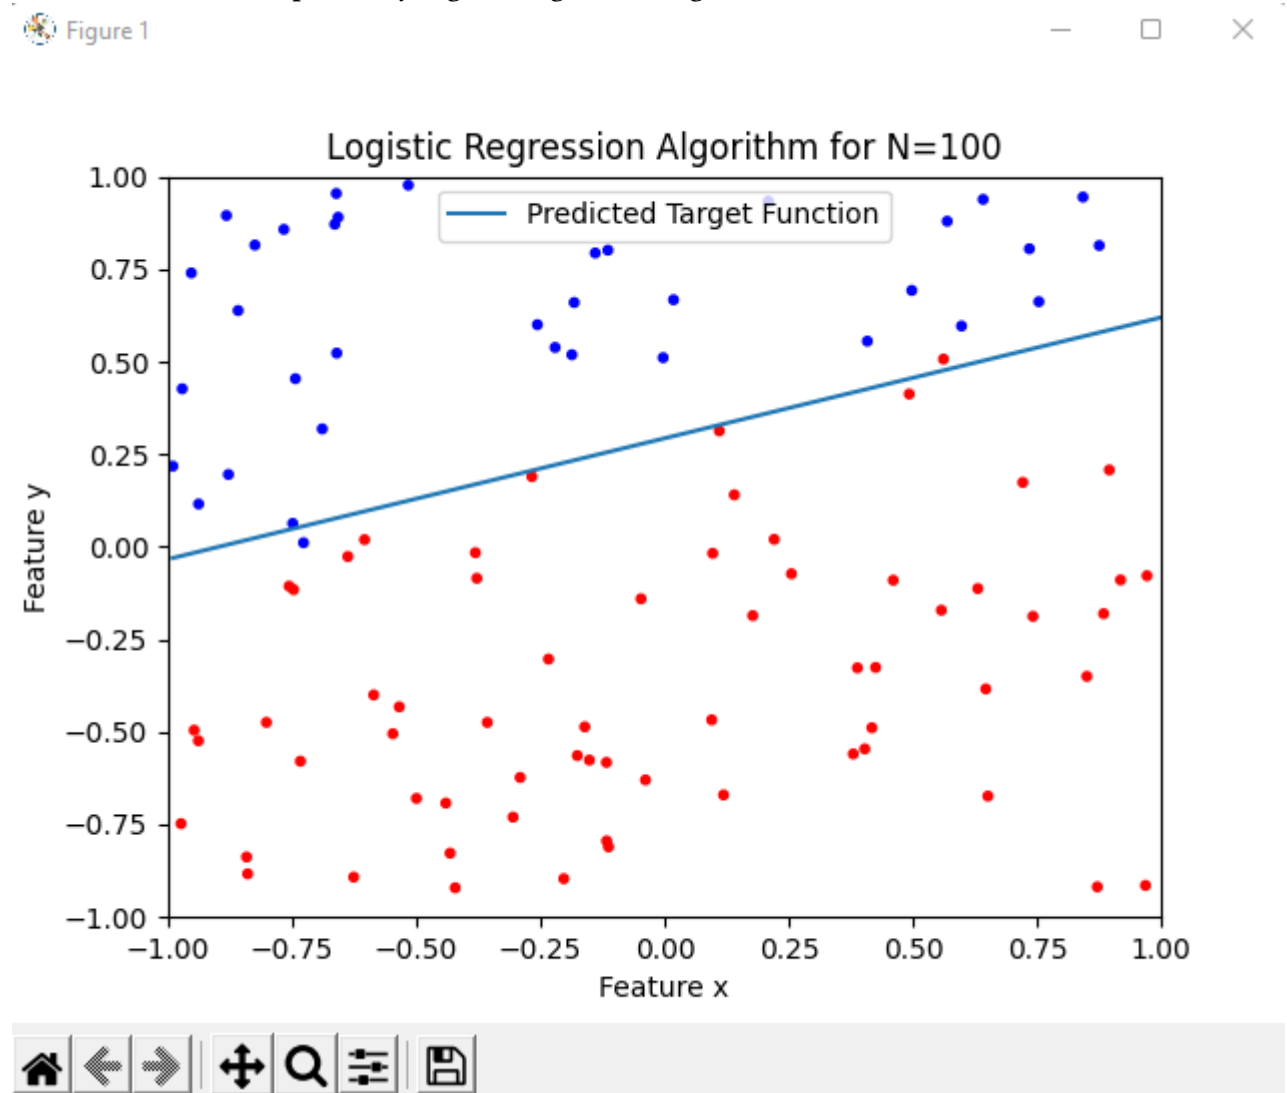
$$\text{Fixed step size: } w(1) = w(0) + n \nabla E_{in}$$

$$v \text{ is calculated by stochastic gradient descent: } \nabla E_{in} = -\frac{1}{N} \sum_{n=1}^N \frac{y_n \cdot x_n}{\ln(1 + e^{y_n \cdot w^T(t) \cdot x_n})}$$

$$\text{cross entropy error: } E_{out}(w) = \frac{1}{N} \sum_{n=1}^N \ln(1 + e^{-y_n \cdot w^T \cdot x_n})$$

More details can be found in the source-code.

Predicted line on 100 points by logistic regression algorithm:



8.

I used the cross entropy error formula to calculate E_{out} .

Average E_{out} after 100 runs: 0.09444365341119473.

The answer is D.

9.

Average epochs after 100 runs: 338.58.

The answer is A.

Regularization with weight decay


For this section, Linear_Regression.py is implemented. Script can be run to observe results. Some data scattering is included.

Results after running Linear_Regression.py:

```
$ py Linear_Regression.py
Q2: Ein: 0.02857142857142857, Eout: 0.084
Q3: Ein: 0.02857142857142857, Eout: 0.08
Q4: Ein: 0.37142857142857144, Eout: 0.436
Q5- Q6: k: -1, Eout: 0.056
```

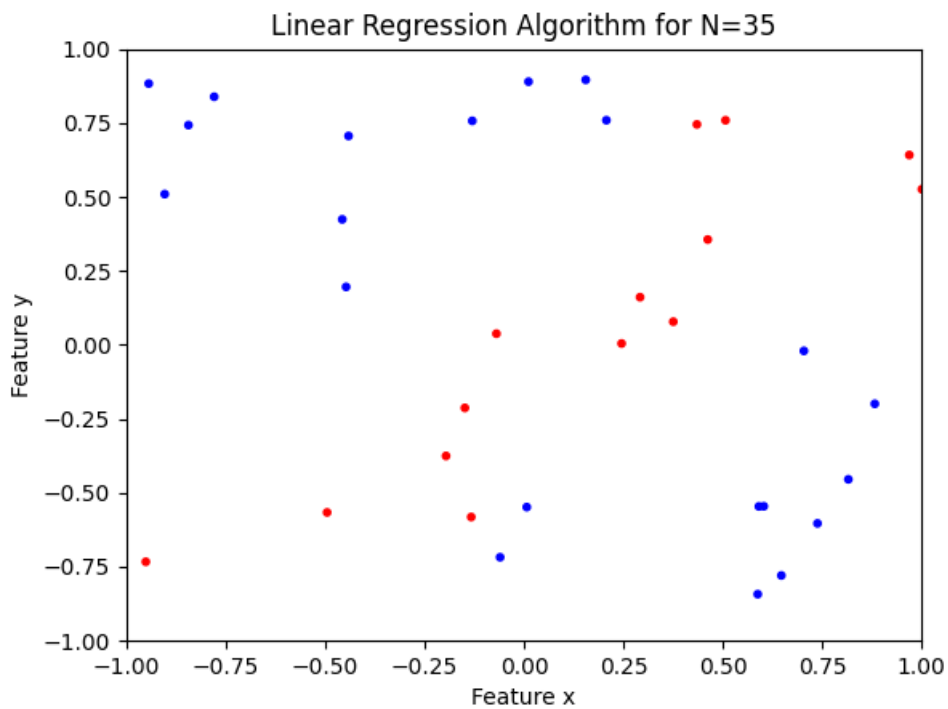
2.

Actual in-sample data:




Predicted in-sample data:

Figure 1

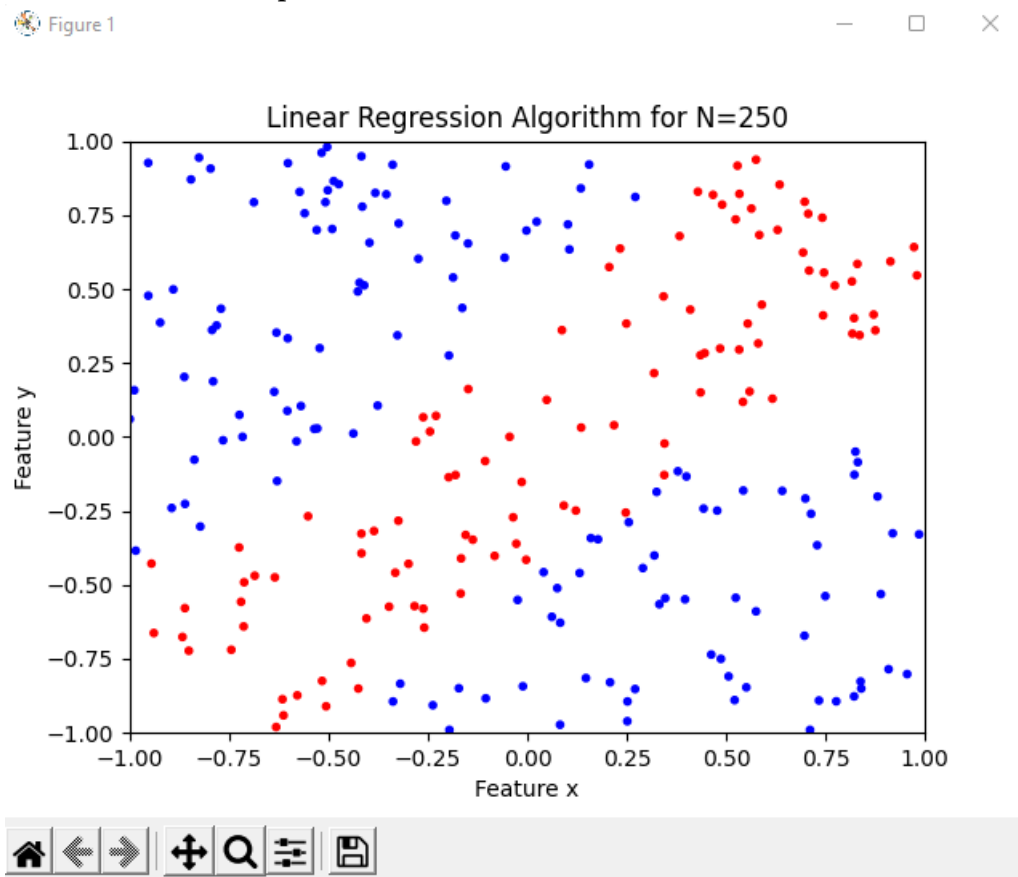


Actual out-of-sample data:

Figure 1



Predicted out-of-sample data:



Ein: 0.02857142857142857, Eout: 0.084

The answer is A.

3.

0.02857142857142857, Eout: 0.08

The answer is D.

4.

Q4: Ein: 0.37142857142857144, Eout: 0.436

The answer is E.

5.

$k = -1$

The answer is D.

6.

Eout: 0.056

The answer is B.

Neural Networks

For this section, Neural_Network.py is implemented using the formulas below. Script can be run to observe results.

Neural Network operators:

$$w_{ij}^{(l)} = \left\{ \begin{array}{ll} 1 \leq l \leq L & \text{Layers} \\ 0 \leq i \leq d^{(l-1)} & \text{Inputs} \\ 1 \leq j \leq d^{(l)} & \text{Outputs} \end{array} \right\}$$

Activation function:

$$\frac{(e^s - e^{-s})}{(e^s + e^{-s})}$$

Forward propagation:

$$x_j^{(l)} = \theta(s_j^{(l)}) = \theta\left(\sum_{i=0}^{d^{(l-1)}} w_{ij}^{(l)} \cdot x_i^{(l-1)}\right)$$

Backward propagation:

$$\text{For final layer: } \delta_1^{(L)} = 2 \cdot (x_1^{(L)} - y_n) \cdot (1 - \theta^2(s_1^{(L)}))$$

$$\text{For other layers: } \delta_i^{(l-1)} = (1 - (x_i^{(l-1)})^2) \cdot \sum_{j=1}^{d^{(l)}} w_{ij}^{(l)} \cdot \delta_j^{(l)}$$

$$\text{Updating perceptron weights: } w_{ij}^{(l)} \leftarrow w_{ij}^{(l)} - n \cdot x_i^{(l-1)} \cdot \delta_j^{(l)}$$

Results after running Neural_Network.py:

```
$ py Neural_Network.py
Q8: Total number of operations is: 28
Q9: Minimum number of weights is: 46
Q10: Maximum number of weights is: 522
```

8.

Total number of operations is: 28. More details can be observed from run_q8() function in the source-code

The answer is A.

9.

By setting 36 hidden layers with 1 weight we count all the weights in the neural network as 46. More details can be observed from run_q9() function in the source-code

The answer is A.

10.

By setting 2 hidden layers with weights of 18, 18 we count all the weights in the neural network as 522. More details can be observed from `run_q10()` function in the source-code.
The answer is E.